


# Why do we need LangSmith

Friday, 15 August 2025

8:05 AM

## AI Job Application Assistant



complex LM workflow

startup → problem

LLM app

10 bar

fetch

matchins

apply ←

resume ← jd  
cover letter

latency → 2-min

slow → 7-10 mins

topic ↴

Receive a query → e.g. "Summarize the top 5 latest papers on solar energy storage."

Search academic sources for relevant papers.

Read each paper and extract key points.

Generate a detailed summary with citations

Answer follow-up questions if the user asks

research assistant

agent

spike inst

debias agent

→ google scholar, arxiv

report → 50 page  
cost/openAI → increase

users → loss

autonomy

reports → 2PS

rank → 50 point

prompt ← user exp



## conversational response

\*  
Observability is the ability to understand a system's internal state by examining its external outputs, like logs, metrics, and traces.

It allows you to diagnose issues, understand performance, and improve reliability by analyzing data generated by the system.

Essentially, it's about being able to answer "why" something is happening within a system, even if you didn't anticipate the problem.

# What is LangSmith

---

Friday, 15 August 2025 8:13 AM

LangSmith is a unified observability & evaluation platform where teams can debug, test, and monitor AI app performance.

# What does LangSmith Trace?

Friday, 25 August 2023 8:11 AM

1. Input and Output ✓
2. All the intermediate steps ✓
3. Latency
4. Token usage
5. Cost
6. Error
7. Tags
8. Metadata
9. Feedback

# Code Examples

I

Friday, 15 August 2025 8:11 AM

RAG apps have two big failure modes:

1. Retriever errors – wrong / irrelevant docs retrieved.
2. Generator errors – model hallucinates or misuses context.

In production, it's often unclear where the *failure happened*. Was the retriever bad, or did the LLM ignore the docs?

- LangSmith automatically records:
  - User query
  - Retrieved documents
  - LLM prompt (with inserted docs)
  - LLM response

- Every graph execution can be logged in LangSmith as a **trace**.
- Each **node** (e.g., retriever, LLM, tool call, subgraph) becomes a **run** inside the trace.
- You can visualize the path taken:
  - START → Retriever → Reranker → LLM Answer → END
- If a workflow branches (conditional/parallel/subgraph), LangSmith captures which path was executed.

## Other Features of LangSmith

Tuesday, 15-August-2023 10:32 AM

57

multiple traces

monitor

### Monitoring and Alerting

#### What it does:

Monitoring in LangSmith looks across many traces at once to track the overall health of your LLM system. It aggregates key metrics like latency (P50, P95, P99), token usage, cost, error rates, and success rates. You can set up alerts to notify you when these metrics drift outside acceptable ranges (e.g., a spike in latency, higher error rates, or unexpected cost growth).

#### Why it matters:

In production, issues often appear first as patterns across multiple runs rather than in a single trace. Monitoring helps you catch these early signals before they impact users at scale. Instead of waiting for customer complaints, you're proactively alerted when performance degrades or costs spike, enabling faster response and more reliable applications.

observe



LLM → exl



trace →

### Evaluation

#### What it does:

Evaluation in LangSmith helps you systematically measure the quality of your LLM outputs. You can run tests against gold-standard datasets or apply custom evaluation metrics such as faithfulness, relevance, or completeness. LangSmith supports multiple approaches: automated scoring with LLM-as-a-judge, semantic similarity checks, or even custom Python evaluators. Evaluations can be run both

## Monitoring and Alerting

### What it does:

Monitoring in LangSmith looks across **many traces at once** to track the overall health of your LLM system. It aggregates key metrics like **latency (P50, P95, P99)**, **token usage**, **cost**, **error rates**, and **success rates**. You can set up **alerts** to notify you when these metrics drift outside acceptable ranges (e.g., a spike in latency, higher error rates, or unexpected cost growth).

### Why it matters:

In production, issues often appear first as **patterns across multiple runs** rather than in a single trace. Monitoring helps you catch these early signals before they impact users at scale. Instead of waiting for customer complaints, you're proactively alerted when performance degrades or costs spike, enabling faster response and more reliable applications.

## Evaluation

### What it does:

Evaluation in LangSmith helps you systematically measure the quality of your LLM outputs. You can run tests against gold-standard datasets or apply custom evaluation metrics such as faithfulness, relevance, or completeness. LangSmith supports multiple approaches: automated scoring with LLM-as-a-judge, semantic similarity checks, or even custom Python evaluators. Evaluations can be run both offline (batch tests before deployment) and online (continuous checks on live traffic).

### Why it matters:

LLM behavior can be unpredictable — a small change in prompts, models, or retrieval logic may improve some cases but break others. Evaluation provides an objective, repeatable way to track performance over time, ensuring that new versions are actually better and preventing regressions.

### Example:

For a RAG chatbot, you might evaluate:

- Faithfulness → Are answers grounded in retrieved documents?
- Relevance → Did the response actually address the user's question?

By running the same dataset across GPT-4, Claude, and LLaMA, you can directly compare which model (or pipeline setup) performs best.

- Relevance → Did the response actually address the user's question?

By running the same dataset across GPT-4, Claude, and LLaMA, you can directly compare which model (or pipeline setup) performs best.

## Prompt Experimentation

What it does:

Prompt Experimentation in LangSmith allows you to systematically test and compare different prompt versions. You can run A/B tests across prompts on the same dataset, track their performance against evaluation metrics, and record the outcomes. Results are stored over time, giving you a clear history of which prompt variations worked best and under what conditions.

## Dataset Creation & Annotation

### What it does:

- Provides tools to build datasets for evaluation and fine-tuning.
- Supports manual annotation (e.g., labeling whether an answer is correct).
- Stores datasets versioned for reuse across projects.

### Why it matters:

High-quality datasets are critical for evaluation and feedback loops.

### Example:

- Customer support: Build a dataset of **common questions + expected answers**.
- Use it to benchmark your RAG agent every time you change retrieval logic.

## User Feedback Integration

What it does:

- Lets you capture thumbs up/down, ratings, or structured feedback from users in production.
- Feedback is logged alongside traces → tied to the exact prompt, model, and state.
- Supports bulk analysis of what users like/dislike.

## Collaboration

What it does:

- Team members can view, share, and comment on traces, datasets, and evaluations.
- Provides a web UI where non-engineers (PMs, QA, annotators) can inspect and annotate runs.
- Enables shared experiment dashboards.